



# GPU programming with OpenACC

October 22 – 23, 2020  
CSC – IT Center for Science Ltd., Espoo

**Martti Louhivuori  
Georgios Markomanolis**



BY SA

All material (C) 2011–2020 by CSC – IT Center for Science Ltd.  
This work is licensed under a **Creative Commons Attribution-ShareAlike**  
4.0 Unported License, <http://creativecommons.org/licenses/by-sa/4.0>

# Introduction to GPUs in HPC



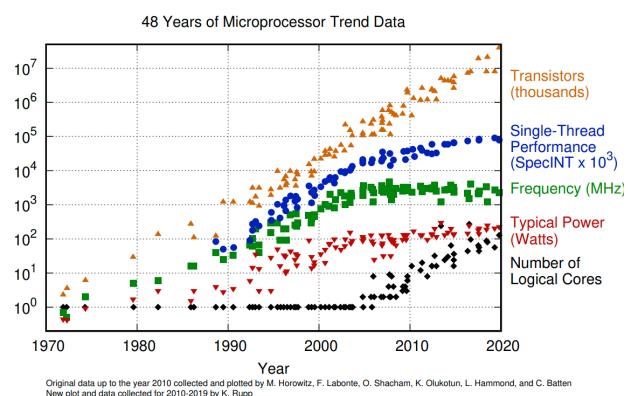
## High-performance computing

- High performance computing is fueled by ever increasing performance
- Increasing performance allows breakthroughs in many major challenges that humankind faces today
- Not only hardware performance, algorithmic improvements have also added ordered of magnitude of real performance



## HPC through the ages

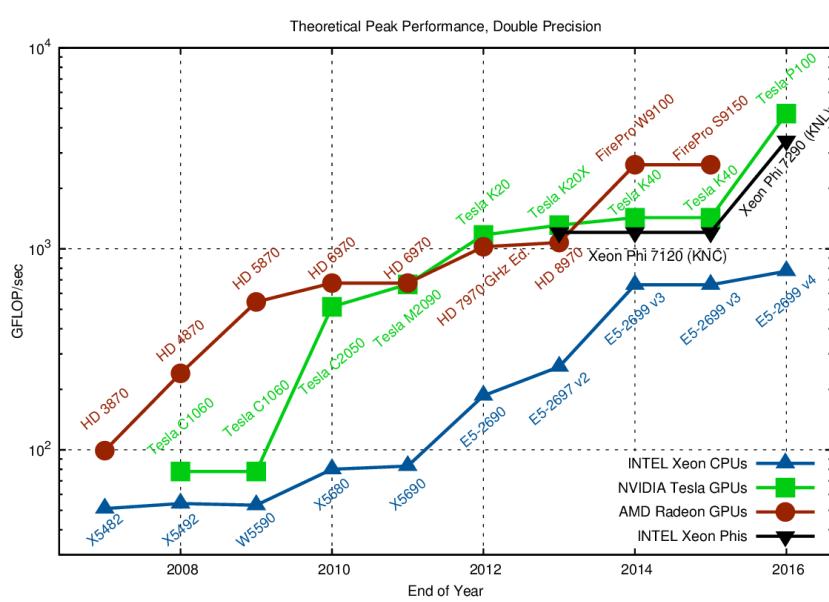
- Achieving performance has been based on various strategies throughout the years
  - Frequency, vectorization, multinode, multicore ...
  - Now performance is mostly limited by power consumption
- Accelerators provide compute resources based on a very high level of parallelism to reach high performance at low relative power consumption



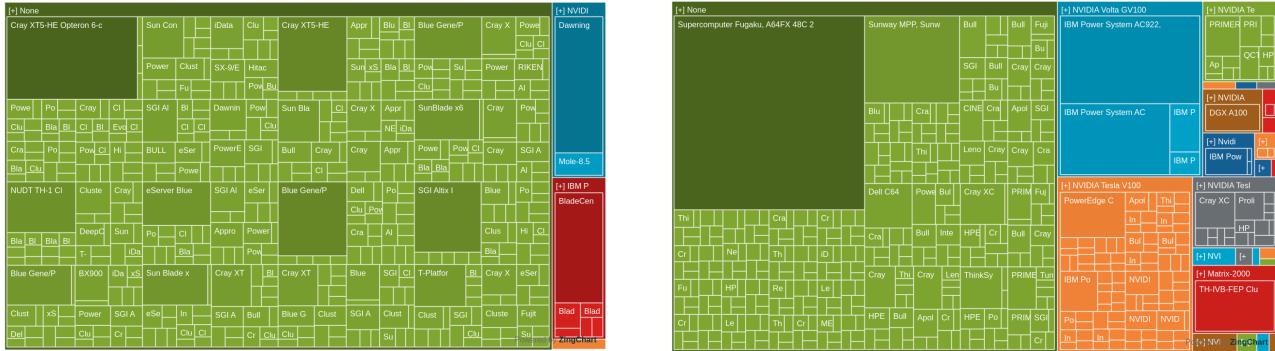
## Accelerators

- Specialized parallel hardware for floating point operations
  - Co-processors for traditional CPUs
  - Based on highly parallel architectures
  - Graphics processing units (GPU) have been the most common accelerators during the last few years
- Promises
  - Very high performance per node
- Usually major rewrites of programs required

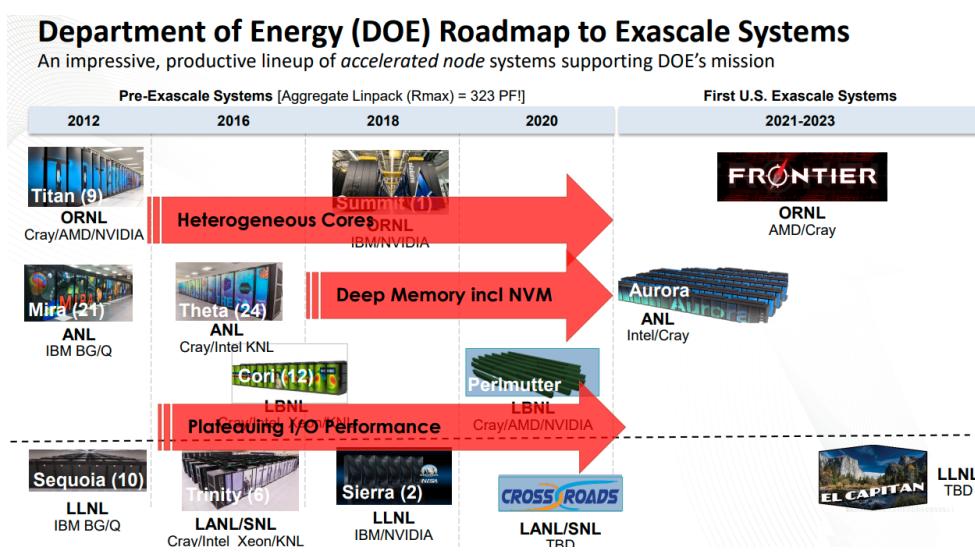
## Accelerator performance growth



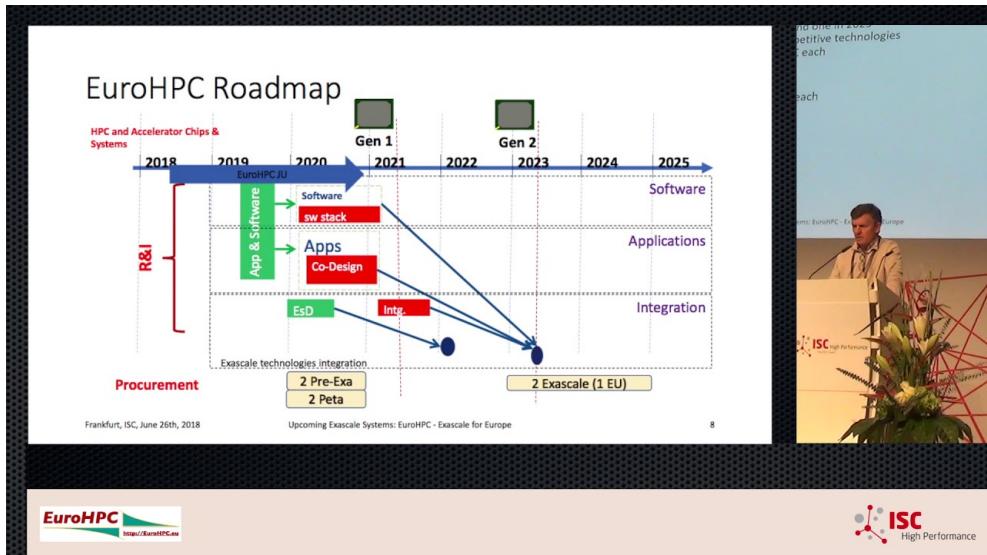
## Accelerators share of 500 fastest systems (Top500) 2010 vs 2020



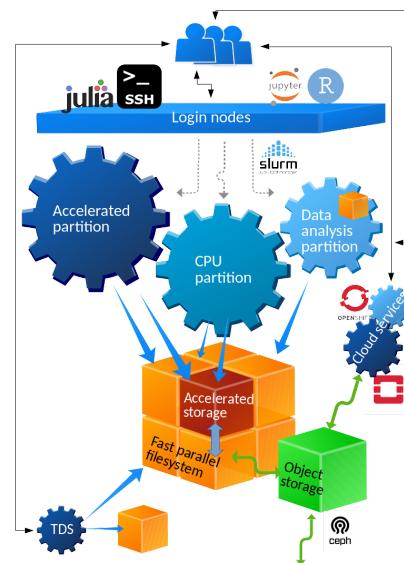
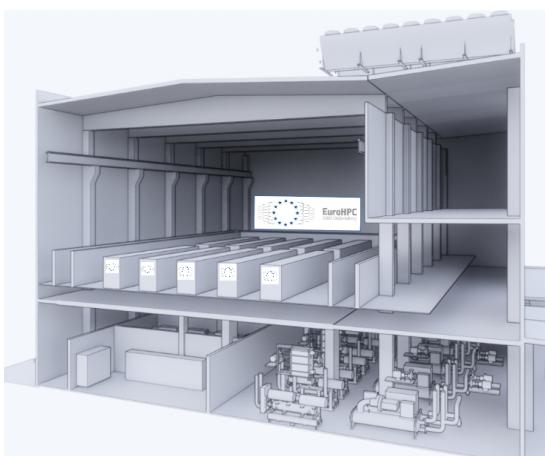
## US roadmap to Exascale



## EU roadmap to Exascale

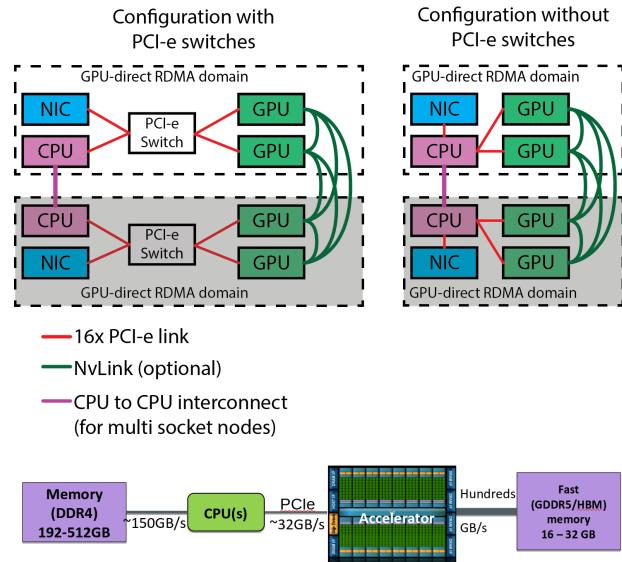


## Lumi - Pre-exascale system in Finland



## Accelerator model today

- Connected to CPUs via PCIe
- Local memory
  - Smaller than main memory (32 GB in Puhti)
  - Very high bandwidth (up to 900 GB/s)
  - Latency high compared to compute performance
- Data must be copied over the PCIe bus



## GPU architecture

- Designed for running tens of thousands of threads simultaneously on thousands of cores
- Very small penalty for switching threads
- Running large amounts of threads hides memory access penalties
- Very expensive to synchronize all threads
- Now Nvidia GPUs have close to monopoly in HPC - will change in next few years

## GPU architecture: Nvidia Volta

- 80 streaming multi processor units (SM), each comprising many smaller Cuda cores
  - 5120 single precision cores
  - 2560 double precision cores
  - 640 tensor cores
- Common L2 cache (6144 KB) for all multi processors
- HBM2 memory, typically 16 GB or 32 GB



## GPU architecture: Nvidia Volta



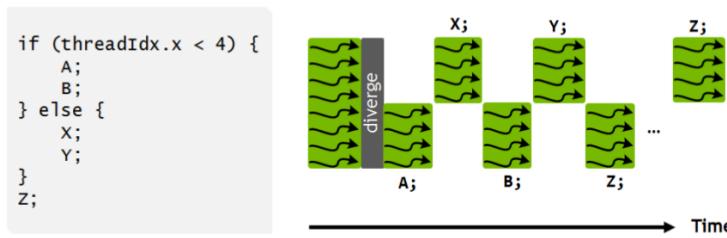
## GPU architecture: Nvidia Volta SM

- 64 single precision cores
- 32 double precision cores
- 64 integer cores
- 8 Tensor cores
- 128 KB memory block for L1 and shared memory
  - 0 - 96 KB can be set to user managed shared memory
  - The rest is L1
- 65536 registers - enables the GPU to run a very large number of threads



## GPU architecture: warps

- All execution is done in terms of 32 threads, a warp
- In a warp 32 threads compute the same instruction on different data (SIMT)
  - Warps are further collected into thread blocks; each executed on one SM
  - In case of divergence (if...) computation is done one branch at a time



## Challenges in using Accelerators

**Applicability:** Is your algorithm suitable for GPU?

**Programmability:** Is the programming effort acceptable?

**Portability:** Rapidly evolving ecosystem and incompatibilities between vendors.

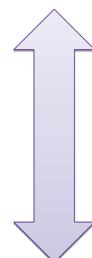
**Availability:** Can you access a (large scale) system with GPUs?

**Scalability:** Can you scale the GPU software efficiently to several nodes?

## Using GPUs

1. Use existing GPU applications
2. Use accelerated libraries
3. Directive based methods
  - OpenMP
  - **OpenACC**
4. Use lower level language
  - CUDA
  - HIP
  - OpenCL

Easier, but more limited



More difficult, but more opportunities

## Directive-based accelerator languages

- Annotating code to pinpoint accelerator-offloadable regions
- OpenACC standard created in Nov 2011
  - Focus on optimizing productivity (reasonably good performance with minimal effort)
  - Current standard is 3.0 (November 2019)
  - Mostly Nvidia only
- OpenMP
  - Earlier only threading for CPUs
  - 4.5 also includes for the first time some support for accelerators
  - 5.0 standard vastly improved
  - Dominant directive approach in the future?

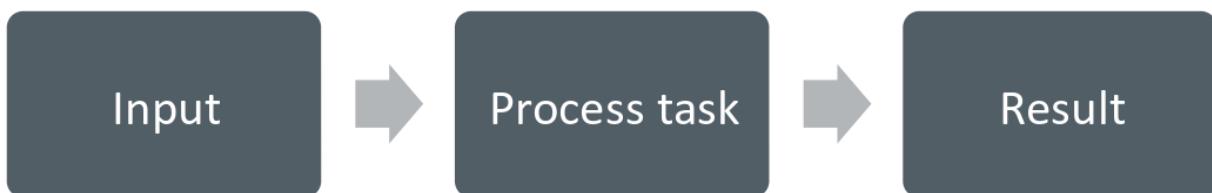
## GPUs at CSC - Puhti-AI

- In total 80 nodes with a total peak performance of 2.7 Petaflops
- Each node has
  - Two latest generation Intel Xeon processors, code name Cascade Lake, with 20 cores each running at 2.1 GHz (Xeon Gold 6230)
  - Four Nvidia Volta V100 GPUs with 32 GB of memory each
  - 384 GB of main memory
  - 3.2 TB of fast local storage
  - Dual rail HDR100 interconnect network connectivity providing 200Gbps aggregate bandwidth

## Parallel computing concepts

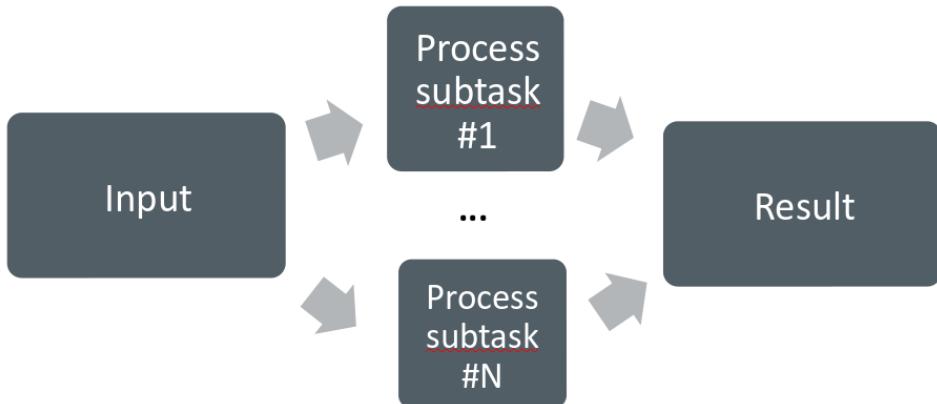
### Computing in parallel

- Serial computing
  - Single processing unit ("core") is used for solving a problem



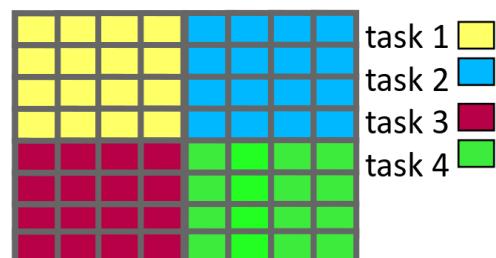
## Computing in parallel

- Parallel computing
  - A problem is split into smaller subtasks
  - Multiple subtasks are processed *simultaneously* using multiple cores



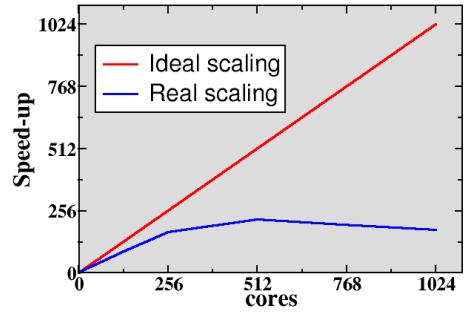
## Exposing parallelism

- Data parallelism
  - Data is distributed to processor cores
  - Each core performs simultaneously (nearly) identical operations with different data
  - Especially good on GPUs(!)
- Task parallelism
  - Different cores perform different operations with (the same or) different data
- These can be combined



## Parallel scaling

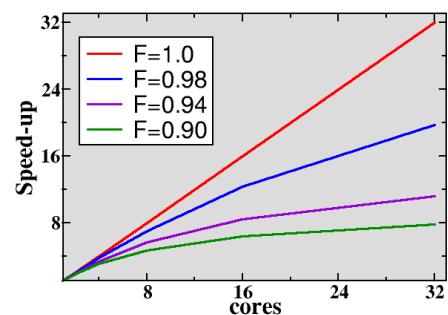
- Strong parallel scaling
  - Constant problem size
  - Execution time decreases in proportion to the increase in the number of cores
- Weak parallel scaling
  - Increasing problem size
  - Execution time remains constant when number of cores increases in proportion to the problem size



## Amdahl's law

- Parallel programs often contain sequential parts
- *Amdahl's law* gives the maximum speed-up in the presence of non-parallelizable parts
- Main reason for limited scaling
- Maximum speed-up is  $\lceil \frac{1}{(1-F) + F/N} \rceil$

where  $F$  is the parallel fraction  
and  $N$  is the number of cores



## Parallel computing concepts

- Load balance
  - Distribution of workload to different cores
- Parallel overhead
  - Additional operations which are not present in serial calculation
  - Synchronization, redundant computations, communications

## Summary

- HPC throughout the ages -- performance through parallelism
- Programming GPUs
  - CUDA, HIP
  - Directive based methods

# OpenACC: introduction



## What is OpenACC ?

- OpenACC defines a set of compiler directives that allow code regions to be offloaded from a host CPU to be computed on a GPU
  - High level GPU programming
  - Large similarity to OpenMP directives
- Supports for both C/C++ and Fortran bindings
- More about OpenACC standard: <http://www.openacc.org>

## OpenACC vs. CUDA/HIP

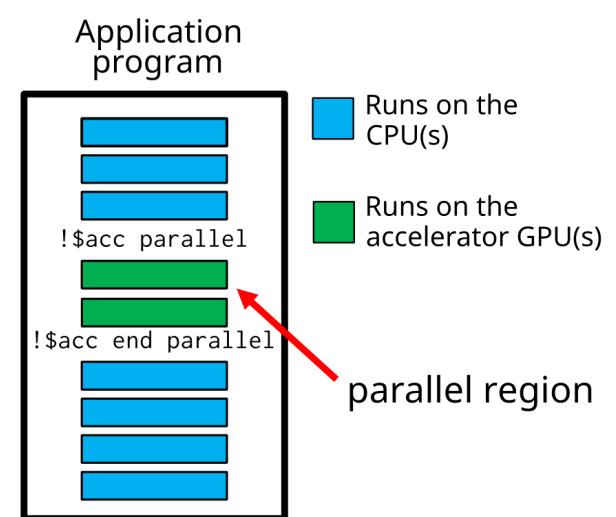
- Why OpenACC and not CUDA/HIP?
  - Easier to work with
  - Porting of existing software requires less work
  - Same code can be compiled to CPU and GPU versions easily
- Why CUDA/HIP and not OpenACC?
  - Can access all features of the GPU hardware
  - More optimization possibilities

## OpenACC execution model

- Host-directed execution with an attached accelerator
  - Large part of the program is usually executed by the host
  - Computationally intensive parts are *offloaded* to the accelerator that executes *parallel regions*
- Accelerator can have a separate memory
  - OpenACC exposes the separate memories through *data environment* that defines the memory management and needed copy operations

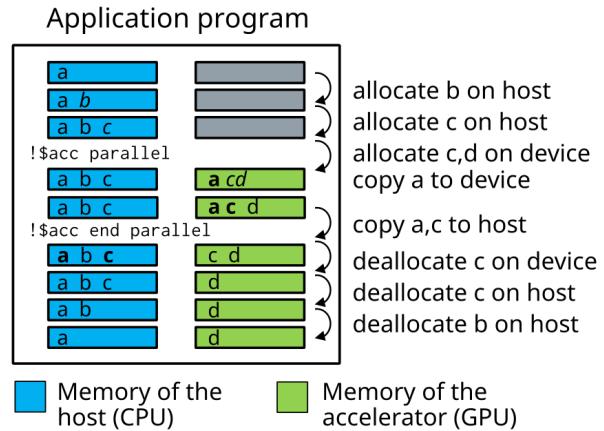
## OpenACC execution model

- Program runs on the host CPU
- Host offloads compute-intensive regions (*kernels*) and related data to the accelerator GPU
- Compute kernels are executed by the GPU



## OpenACC data model

- If host memory is separate from accelerator device memory
  - host manages memory of the device
  - host copies data to/from the device
- When memories are not separate, no copies are needed (difference is transparent to the user)



## OpenACC directive syntax

	<b>sentinel</b>	<b>construct</b>	<b>clauses</b>
C/C++	#pragma acc kernels		copy(data)
Fortran	!\$acc kernels		copy(data)

- OpenACC uses compiler directives for defining compute regions (and data transfers) that are to be performed on a GPU
- Important constructs
  - parallel, kernels, data, loop, update, host\_data, wait
- Often used clauses
  - if (condition), async(handle)

## Compiling an OpenACC program

- Compilers that support OpenACC usually require an option that enables the feature
  - PGI: -acc
  - Cray: -h acc
  - GNU (partial support): -fopenacc
- Without these options a regular CPU version is compiled!

## OpenACC conditional compilation

- Conditional compilation with \_OPENACC macro:

```
#ifdef _OPENACC
device specific code
#else
host code
#endif
```

- \_OPENACC macro is defined as *yyyymm*, where *yyy* and *mm* refer to the year and month of when the specification supported by the compiler was released



## OpenACC internal control variables

- OpenACC has internal control variables
  - ACC\_DEVICE\_TYPE controls which type of accelerator device is to be used.
  - ACC\_DEVICE\_NUM controls which accelerator of the selected type is used.
- During runtime, values can be modified or queried with  
`acc_<set|get>_device_<type|num>`
- Values are always re-read before a kernel is launched and can be different for different kernels



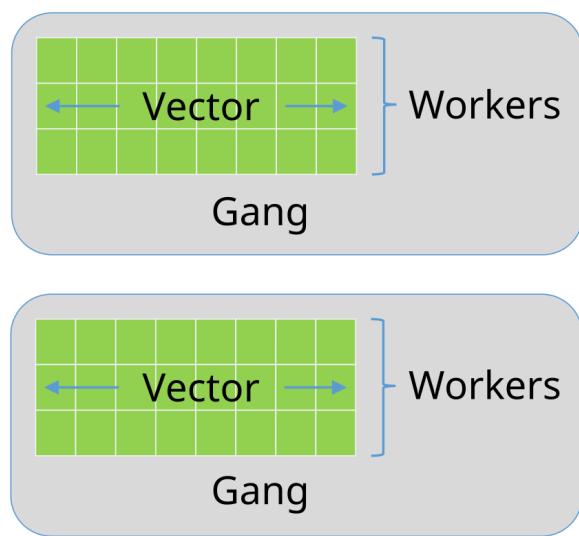
## Runtime API functions

- Low-level runtime API functions can be used to
  - Query the number and type of devices in the system
  - Initialize/shutdown the device(s)
  - Allocate/deallocate memory on the device(s)
  - Transfer data to/from the device(s)
- Function definitions are in
  - C/C++ header file `openacc.h`
  - `openacc` Fortran module (`openacc_lib.h` header in some implementations)

## OpenACC compute constructs

### OpenACC levels of parallelism

- OpenACC has three levels of parallelism
  - **Vector** threads work in SIMD (SIMT) fashion
  - **Workers** compute a vector
  - **Gangs** have one or more workers that share resources, such as streaming multiprocessor
  - Multiple gangs work independently



## OpenACC compute constructs

- OpenACC includes two different approaches for defining parallel regions
  - `parallel` defines a region to be executed on an accelerator. Work sharing *parallelism* has to be defined *manually*. Good tuning prospects.
  - `kernels` defines a region to be transferred into a series of kernels to be executed in *sequence* on an accelerator. Work sharing parallelism is defined *automatically* for the separate kernels, but tuning prospects limited.
- With similar work sharing, both can perform equally well

## Compute constructs: kernels

- Define a region to be transferred to a sequence of kernels for execution on the accelerator device
  - C/C++: `#pragma acc kernels [clauses]`
  - Fortran: `!$acc kernels [clauses]`
- Each separate *loop nest* inside the region will be converted into a separate *parallel kernel*
- The *kernel*s will be executed in a *sequential* order

## Example: kernels

### C/C++

```
/* Compute y=a*x+y */
void accdaxpy(int n, double a,
              const double * restrict x,
              double * restrict y)
{
    #pragma acc kernels
    for (int j=0; j<n; ++j)
        y[j] += a * x[j];
}

...
/* An example call to accdaxpy */
accdaxpy(1<<16, 3.14, x, y);
```

### Fortran

```
! Compute y=a*x+y
subroutine accdaxpy(n, a, x, y)
integer :: n, j
real(kind=8) :: a, x(n), y(n)

 !$acc kernels
do j = 1,n
    y(j) = y(j) + a * x(j)
end do
 !$acc end kernels
end subroutine accdaxpy

...
! An example call to accdaxpy
call accdaxpy(65536, 3.14D0, x, y)
```

## Compute constructs: parallel

- Define a region to be executed on the accelerator device
  - C/C++: #pragma acc parallel [clauses]
  - Fortran: !\$acc parallel [clauses]
- Without any *work sharing* constructs, the whole region is executed *redundantly* multiple times
  - Given a sequence of loop nests, each loop nest may be executed simultaneously

## Work sharing construct: loop

- Define a loop to be parallelized
  - C/C++: #pragma acc loop [clauses]
  - Fortran: !\$acc loop [clauses]
  - Must be followed by a C/C++ or Fortran loop construct.
  - Combined constructs with parallel and kernels
    - #pragma acc kernels loop / !\$acc kernels loop
    - #pragma acc parallel loop / !\$acc parallel loop
- Similar in functionality to OpenMP for/do construct
- Loop index variables are private variables by default

## Example: parallel

### C/C++

```
/* Compute y=a*x+y */
void accdaxpy(int n, double a,
              const double * restrict x,
              double * restrict y)
{
    #pragma acc parallel loop
    for (int j=0; j<n; ++j)
        y[j] += a * x[j];
}

...
/* An example call to accdaxpy */
accdaxpy(1<<16, 3.14, x, y);
```

### Fortran

```
! Compute y=a*x+y
subroutine accdaxpy(n, a, x, y)
    integer :: n, j
    real(kind=8) :: a, x(n), y(n)

    !$acc parallel loop
    do j = 1,n
        y(j) = y(j) + a * x(j)
    end do
    !$acc end parallel loop
end subroutine accdaxpy

...
! An example call to accdaxpy
call accdaxpy(65536, 3.14D0, x, y)
```

## Compiler diagnostics

## Compiler diagnostics

- Compiler diagnostics is usually the first thing to check when starting the OpenACC work
  - It can tell you what operations were actually performed
  - Data copies that were made
  - If and how the loops were parallelized
- The diagnostics is very compiler dependent
  - Compiler flags
  - Level and formatting of information



## PGI compiler

- Diagnostics is controlled by compiler flag -Minfo=option
- Useful options:
  - accel -- operations related to the accelerator
  - all -- print all compiler output
  - intensity -- print loop computational intensity info
  - ccff -- add extra information to the object files for use by tools



## Example: -Minfo

```
$ pgcc -fast -Minfo=all -c util.c
malloc_2d:
    28, Loop not vectorized: data dependency
        Loop unrolled 8 times
        Generated 1 prefetches in scalar loop
eval_point:
    38, Loop not vectorized/parallelized: potential early exits

$ pgcc -fast -Minfo=intensity -c util.c
malloc_2d:
    28, Intensity = 3.00
eval_point:
    38, Intensity = 8.00
```



## Example: -Minfo

```
$ pgcc -acc -Minfo=all doubleloops.c
init:
 38, Memory zero idiom, loop replaced by call to __c_mzero8
 44, Memory set idiom, loop replaced by call to __c_mset8
main:
 74, Generating Tesla code
    77, #pragma acc loop gang /* blockIdx.x */
    79, #pragma acc loop vector(128) /* threadIdx.x */
 74, Generating implicit copyin(u[:1024][:1024]) [if not already present]
    Generating implicit copyout(unew[1:1022][1:1022]) [if not already present]
 79, Loop is parallelizable
 84, Generating Tesla code
    87, #pragma acc loop gang /* blockIdx.x */
    89, #pragma acc loop vector(128) /* threadIdx.x */
 84, Generating implicit copyin(unew[:1024][:1024]) [if not already present]
    Generating implicit copyout(u[1:1022][1:1022]) [if not already present]
 89, Loop is parallelizable
```



## Summary

- OpenACC is an directive-based extension to C/Fortran programming languages for accelerators
- Supports separate memory on the accelerator
- Compute constructs: parallel and kernels
- Compiler diagnostics

# OpenACC: data management



## OpenACC data environment

- OpenACC supports devices which either share memory with or have a separate memory from the host
- Constructs and clauses for
  - defining the variables on the device
  - transferring data to/from the device
- All variables used inside the parallel or kernels region will be treated as *implicit* variables if they are not present in any data clauses, i.e. copying to and from the device is automatically performed

## Motivation for optimizing data movement

- When dealing with an accelerator / GPU device attached to a PCIe bus, **optimizing data movement** is often **essential** to achieving good performance
- The four key steps in porting to high performance accelerated code:
  1. Identify parallelism
  2. Express parallelism
  3. Express data movement
  4. Optimise loop performance
  5. Go back to 1!

## Data lifetimes

- Typically data on the device has the same lifetime as the OpenACC construct (parallel, kernels, data) it is declared in
- It is possible to declare and refer to data residing statically on the device until deallocation takes place
- If the accelerator has a separate memory from the host, any modifications to the data on the host are not visible to the device before an explicit update has been made and vice versa

## Data constructs: data

- Define a region with data declared in the device memory
  - C/C++: #pragma acc data [clauses]
  - Fortran: !\$acc data [clauses]
- Data transfers take place
  - from the **host** to the **device** upon entry to the region
  - from the **device** to the **host** upon exit from the region
- Functionality defined by *data clauses*
- *Data clauses* can also be used in kernels and parallel constructs

## Data construct: example

### C/C++

```
float a[100];
int iter;
int maxit=100;

#pragma acc data create(a)
{
    /* Initialize data on device */
    init(a);
    for (iter=0; iter < maxit; iter++) {
        /* Computations on device */
        acc_compute(a);
    }
}
```

### Fortran

```
real :: a(100)
integer :: iter
integer, parameter :: maxit = 100

!$acc data create(a)
! Initialize data on device
call init(a)
do iter=1,maxit
    ! Computations on device
    call acc_compute(a)
end do
!$acc end data
```

## Data constructs: data clauses

### present(var-list)

- **on entry/exit:** assume that memory is allocated and that data is present on the device

### create(var-list)

- **on entry:** allocate memory on the device, unless it was already present
- **on exit:** deallocate memory on the device, if it was allocated on entry
  - in-depth: *structured* reference count decremented, and deallocation happens if both reference counts (*structured* and *dynamic*) are zero



## Data constructs: data clauses

`copy(var-list)`

- **on entry:** if data is present on the device on entry, behave as with the present clause, otherwise allocate memory on the device and copy data from the host to the device.
- **on exit:** copy data from the device to the host and deallocate memory on the device if it was allocated on entry



## Data constructs: data clauses

`copyin(var-list)`

- **on entry:** same as copy on entry
- **on exit:** deallocate memory on the device if it was allocated on entry

`copyout(var-list)`

- **on entry:** if data is present on the device on entry, behave as with the present clause, otherwise allocate memory on the device
- **on exit:** same as copy on exit

## Data constructs: data clauses

```
reduction(operator:var-list)
```

- Performs reduction on the (scalar) variables in list
- Private reduction variable is created for each gang's partial result
  - initialised to operators initial value
- After parallel region the reduction operation is applied to the private variables and the result is aggregated to the shared variable *and* the aggregated result is combined with the original value of the variable

## Reduction operators in C/C++ and Fortran

Arithmetic Operator	Initial value
+	0
-	0
*	1
max	least
min	largest

## Reduction operators in C/C++ only

Logical Operator	Initial value
<code>&amp;&amp;</code>	1
<code>  </code>	0

Bitwise Operator	Initial value
<code>&amp;</code>	<code>~0</code>
<code> </code>	0
<code>^</code>	0

## Reduction operators in Fortran

Logical Operator	Initial value
<code>.and.</code>	<code>.true.</code>
<code>.or.</code>	<code>.false.</code>
<code>.eqv.</code>	<code>.true.</code>
<code>.neqv.</code>	<code>.false.</code>

Bitwise Operator	Initial value
<code>iand</code>	all bits on
<code>ior</code>	0
<code>ieor</code>	0

## Data specification

- Data clauses specify functionality for different variables
- Overlapping data specifications are not allowed
- For array data, *array ranges* can be specified
  - C/C++: arr[start\_index:length], for instance vec[0:n]
  - Fortran: arr(start\_index:end\_index), for instance vec(1:n)
- Note: array data **must** be *contiguous* in memory (vectors, multidimensional arrays etc.)

## Default data environment in compute constructs

- All variables used inside the parallel or kernels region will be treated as *implicit* variables if they are not present in any data clauses, i.e. copying to and from the device is automatically performed
- Implicit *array* variables are treated as having the copy clause in both cases
- Implicit *scalar* variables are treated as having the
  - copy clause in kernels
  - firstprivate clause in parallel

## data construct: example

### C/C++

```
int a[100], d[3][3], i, j;

#pragma acc data copy(a[0:100])
{
    #pragma acc parallel loop present(a)
    for (int i=0; i<100; i++)
        a[i] = a[i] + 1;
    #pragma acc parallel loop \
        collapse(2) copyout(d)
    for (int i=0; i<3; ++i)
        for (int j=0; j<3; ++j)
            d[i][j] = i*3 + j + 1;
}
```

### Fortran

```
integer a(0:99), d(3,3), i, j

!$acc data copy(a(0:99))
 !$acc parallel loop present(a)
 do i=0,99
     a(i) = a(i) + 1
 end do
 !$acc end parallel loop
 !$acc parallel loop collapse(2) copyout(d)
 do j=1,3
     do i=1,3
         d(i,j) = i*3 + j + 1
     end do
 end do
 !$acc end parallel loop
 !$acc end data
```

## Unstructured data regions

- Unstructured data regions enable one to handle cases where allocation and freeing is done in a different scope
- Useful for e.g. C++ classes, Fortran modules
- `enter data` defines the start of an unstructured data region
  - C/C++: `#pragma acc enter data [clauses]`
  - Fortran: `!$acc enter data [clauses]`
- `exit data` defines the end of an unstructured data region
  - C/C++: `#pragma acc exit data [clauses]`
  - Fortran: `!$acc exit data [clauses]`



## Unstructured data

```
class Vector {  
    Vector(int n) : len(n) {  
        v = new double[len];  
        #pragma acc enter data create(v[0:len])  
    }  
    ~Vector() {  
        #pragma acc exit data delete(v[0:len])  
        delete[] v;  
    }  
    double v;  
    int len;  
};
```



## Enter data clauses

`if(condition)`

- Do nothing if condition is false

`create(var-list)`

- Allocate memory on the device

`copyin(var-list)`

- Allocate memory on the device and copy data from the host to the device

## Exit data clauses

`if(condition)`

- Do nothing if condition is false

`delete(var-list)`

- Deallocate memory on the device
  - in-depth: *dynamic* reference count decremented, and deallocation happens if both reference counts (*dynamic* and *structured*) are zero

`copyout(var-list)`

- Deallocate memory on the device and copy data from the device to the host
  - in-depth: *dynamic* reference count decremented, and deallocation happens if both reference counts (*dynamic* and *structured*) are zero

## Data directive: update

- Define variables to be updated within a data region between host and device memory
  - C/C++: `#pragma acc update [clauses]`
  - Fortran: `!$acc update [clauses]`
- Data transfer direction controlled by `host(var-list)` or `device(var-list)` clauses
  - `self (host)` clause updates variables from device to host
  - `device` clause updates variables from host to device
- At least one data direction clause must be present

## Data directive: update

- update is a single line executable directive
- Useful for producing snapshots of the device variables on the host or for updating variables on the device
  - Pass variables to host for visualization
  - Communication with other devices on other computing nodes
- Often used in conjunction with
  - Asynchronous execution of OpenACC constructs
  - Unstructured data regions

## update directive: example

### C/C++

```
float a[100];
int iter;
int maxit=100;

#pragma acc data create(a) {
    /* Initialize data on device */
    init(a);
    for (iter=0; iter < maxit; iter++) {
        /* Computations on device */
        acc_compute(a);
        #pragma acc update self(a) \
                    if(iter % 10 == 0)
    }
}
```

### Fortran

```
real :: a(100)
integer :: iter
integer, parameter :: maxit = 100

!$acc data create(a)
    ! Initialize data on device
    call init(a)
do iter=1,maxit
    ! Computations on device
    call acc_compute(a)
    !$acc update self(a)
    !$acc& if(mod(iter,10)==0)
end do
!$acc end data
```

## Data directive: declare

- Makes a variable resident in accelerator memory
- Added at the declaration of a variable
- Data life-time on device is the implicit life-time of the variable
  - C/C++: #pragma acc declare [clauses]
  - Fortran: !\$acc declare [clauses]
- Supports usual data clauses, and additionally
  - device\_resident
  - link

## Porting and managed memory

- Porting a code with complicated data structures can be challenging because every field in type has to be copied explicitly
- Recent GPUs have *Unified Memory* and support for page faults

```
typedef struct points {
    double x, y;
    int n;
}

void init_point() {
    points p;

    #pragma acc data create(p)
    {
        p.size = n;
        p.x = (double)malloc(...);
        p.y = (double)malloc(...);
        #pragma acc update device(p)
        #pragma acc copyin (p.x[0:n]...)
```



## Managed memory

- Managed memory copies can be enabled on PGI compilers
  - Pascal (P100): --ta=tesla,cc60,managed
  - Volta (V100): --ta=tesla,cc70,managed
- For full benefits Pascal or Volta generation GPU is needed
- Performance depends on the memory access patterns
  - For some cases performance is comparable with explicitly tuned versions



## Summary

- Data directive
  - Structured data region
  - Clauses: copy, present, copyin, copyout, create
- Enter data & exit data
  - Unstructured data region
- Update directive
- Declare directive

# Profiling and performance optimisation





## Outline

- NVidia profiling tools
- Tuning compute region performance
  - Enabling vectorization
  - Loop optimizations
  - Branches
  - Memory access



## Profiling tools

## NVIDIA NVPROF profiler

- NVIDIA released the Nsight recently for profiling. One of the main reasons for the new tool is scalability and, of course, new features
  - It is included with CUDA since 10.x
- GPU profiling capabilities
  - High-level usage statistics
  - Timeline collection
  - Analysis metrics

## OpenACC example

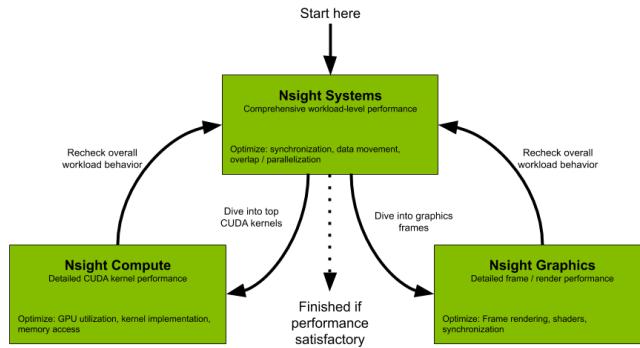
```
$ nsys profile -t nvtv,openacc --stats=true -s cpu ./jacobi
...
Generating CUDA API Statistics...
CUDA API Statistics (nanoseconds)

Time(%)      Total Time      Calls      Average      Minimum      Maximum      Name
-----  -----
  77.1      108005341      1450      74486.4      1815      189044      cuStreamSync
  16.6       23232332       1       23232332.0      23232332      23232332      cuMemHostAlloc
   2.9       4037274       1091      3700.5       2815      24189       cuLaunchKernel
   0.8       1182112       361      3274.5       2911      21123      cuMemcpyDtoH
   0.6        855308       361      2369.3       2065      10502      cuMemsetD32As
   0.6        813341       1       813341.0      813341      813341      cuMemAllocHost

Generating CUDA Kernel Statistics...
CUDA Kernel Statistics (nanoseconds)

Time(%)      Total Time      Instances      Average      Minimum      Maximum      Name
-----  -----
  43.2       43864453       361      121508.2      118208      124671      update_65_gpu
```

# NVIDIA Nsight workflow



Source: NVIDIA

# NVIDIA Systems





## NVIDIA Metrics

```
srun -n 1 nv-nvsiight-cu-cli --devices 0 --query-metrics > my_metrics.txt

dram_bytes          # of bytes accessed in DRAM
dram_bytes_read    # of bytes read from DRAM
dram_bytes_write   # of bytes written to DRAM
dram_cycles_active # of cycles where DRAM was active
...
tpc_cycles_active  # of cycles where TPC was active
tpc_cycles_elapsed # of cycles where TPC was active
tpc_cycles_in_frame # of cycles in user-defined frame
tpc_cycles_in_region # of cycles in user-defined region
...
```



## Nsight Compute CLI (I)

```
srun -n 1 nv-nvsiight-cu-cli ./jacobi
...
update_65_gpu, 2020-Oct-18 23:42:06, Context 1, Stream 13
Section: GPU Speed Of Light
-----
DRAM Frequency           cycle/usecond
SM Frequency            cycle/nsecond
Elapsed Cycles          cycle
Memory [%]              %
SOL DRAM                 %
Duration                  usecond
SOL L1/TEX Cache         %
SOL L2 Cache              %
SM Active Cycles          cycle
SM [%]                   %

WRN Memory is more heavily utilized than Compute: Look at the Memory Workload Analysis report
where the memory system bottleneck is. Check memory replay (coalescing) metrics to make
```

## Nsight Compute CLI (II)

Section: Launch Statistics		
Block Size		128
Grid Size		45,000
Registers Per Thread	register/thread	30
Shared Memory Configuration Size	Kbyte	16.38
Driver Shared Memory Per Block	byte/block	0
Dynamic Shared Memory Per Block	Kbyte/block	1.02
Static Shared Memory Per Block	byte/block	0
Threads	thread	5,760,000
Waves Per SM		35.16

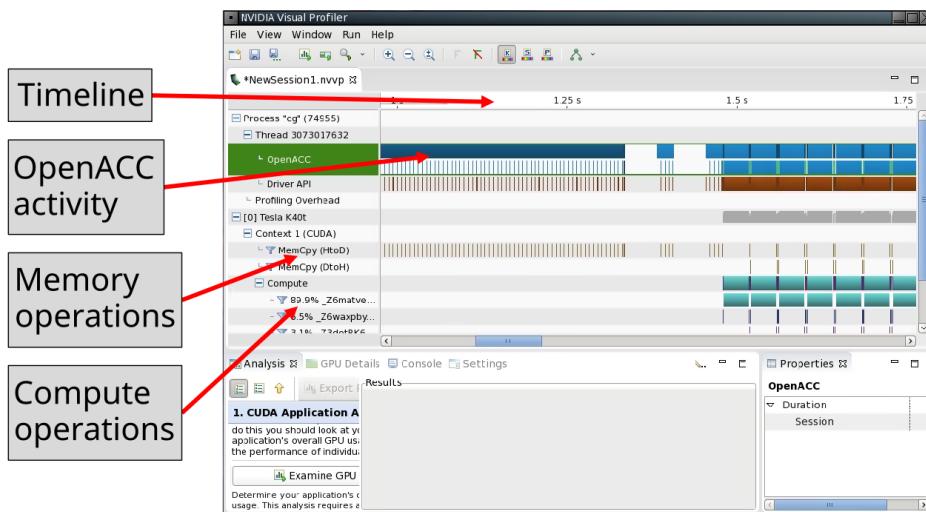
  

Section: Occupancy		
Block Limit SM	block	32
Block Limit Registers	block	16
Block Limit Shared Mem	block	96
Block Limit Warps	block	16
Theoretical Active Warps per SM	warp	64
Theoretical Occupancy	%	100
Achieved Occupancy	%	83.71
Achieved Active Warps Per SM	warp	53.58

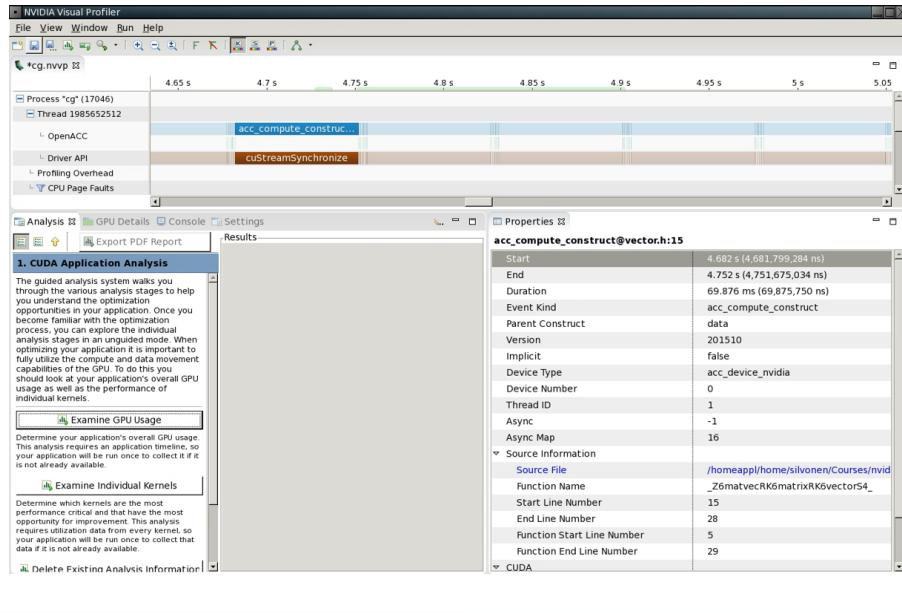
- Nsight Compute GUI does not support OpenACC

## NVIDIA visual profiler

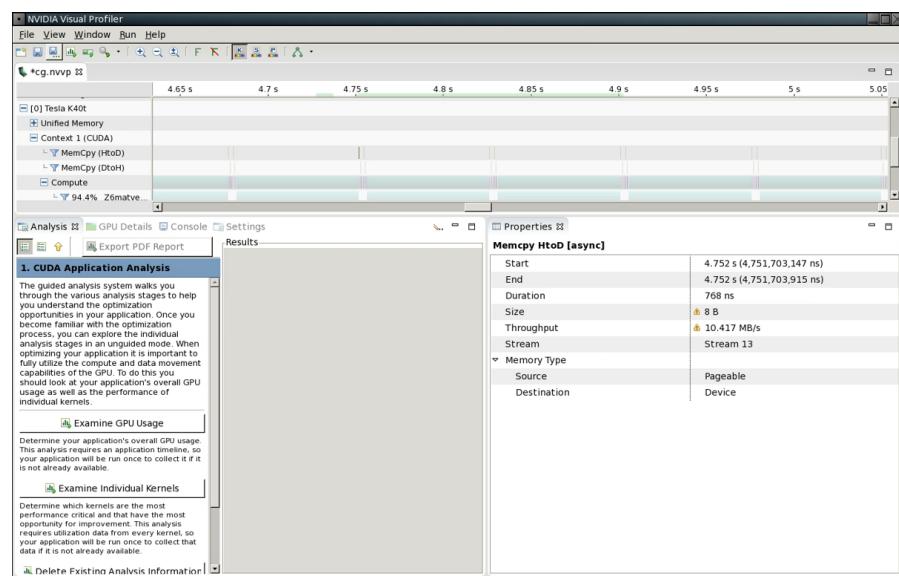
- Nvprof is an older profiling tool



## Details on OpenACC compute construct



## Details on memory copy





## Optimization



### Compute optimizations

- Data movement is an important part to optimize when using GPUs
  - Keeping data on the GPU as long as possible
- Getting the compiler to generate parallel code
  - Addressing loop dependencies
- Data access and execution divergence are important for GPU performance

## Loop dependencies

```
/* FLOW dependency, k>0 */
for (int i=0; i<N; i++)
    A[i] = A[i-k]+1;
/* ANTI dependency, k>0 */
for (int i=0; i<N; i++)
    A[i] = A[i+k]+1;
```

```
! FLOW dependency, k>0
do i=0, N
    a(i) = a(i-k) + 1;
end do
! ANTI dependency, k>0
do i=0, N
    a(i) = a(i+k) + 1;
end do
```

- FLOW dependency
  - Read After Write (RAW), data is written to is read on the following iteration round(s)
- ANTI dependency
  - Write After Read (WAR), data read is written to on the following iteration rounds

## Loop dependencies

- Dependencies disable vectorization, which is essential for good performance
- Rewrite the loops so that the dependency is removed
- Try to split the loop, use temporary array, etc.
- Some dependencies can not be removed
  - Try a different algorithm?

## Loop dependencies and C

- C pointers are hard for the compiler to follow
  - Compiler will not know, if a loop can be vectorized safely, if a function has pointer arguments
  - Can be a *false* dependency

```
void adder(float *x, float *y, float *res) {  
    for (int i=0; i < VECSIZE; i++) {  
        res[i] = x[i] + y[i];  
    }  
}
```

- What if `res` and `x` overlap in memory?

## C99 restrict keyword

- C99 standard has `restrict` keyword which tells the compiler that the pointer is accessed so that it does not overlap with other accesses

```
void adder(float restrict *x, float restrict *y, float restrict *res) {  
    for (int i=0; i < VECSIZE; i++) {  
        res[i] = x[i] + y[i];  
    }  
}
```

## Loop independent clause

- OpenACC independent clause tells to the compiler that loop iterations are independent
  - Overrides any compiler dependency analysis
  - You have to make sure that the iterations are independent!

```
#pragma acc loop independent
void adder(float *x, float *y, float *res) {
    for (int i=0; i < VECSIZE; i++) {
        res[i] = x[i] + y[i];
    }
}
```

## Loop directive

- Loop directive accepts several fine-tuning clauses
  - **gang** -- apply gang-level parallelism
  - **worker** -- apply worker-level parallelism
  - **vector** -- apply vector-level parallelism
  - **seq** -- run sequentially
- Multiple levels can be applied to a loop nest, but they have to be applied in top-down order

## Optimize loops: vector length

- Tell the compiler that when using NVIDIA device it should use a vector length of 32 on the innermost loop
- Because these parameters depend on the accelerator type, it is a good practice to add **device\_type** clause

```
for (int i=0; i<imax; i++) {  
    ...  
    #pragma acc loop device_type(nvidia) vector(32)  
    for (int j=0; j<jmax; j++) {  
        ... /* No further loops in this block */  
    }  
}
```

## Optimize loops: specifying workers

```
#pragma acc loop device_type(nvidia) gang worker(32)  
for (int i=0; i<imax; i++) {  
    ...  
    #pragma acc loop device_type(nvidia) vector(32)  
    for (int j=0; j<jmax; j++) {  
        ...  
    }  
}
```

- Tell the compiler that when using NVIDIA device, the outer loop should be broken over gangs and workers with 32 workers per gang

## Additional loop optimizations

- collapse(N)
  - Same as in OpenMP, take the next N tightly nested loops and flatten them into a one loop
  - Can be beneficial when loops are small
  - Breaks the next loops into tiles (blocks) before parallelizing the loops
  - For certain memory access patterns this can improve data locality

## What values should I try?

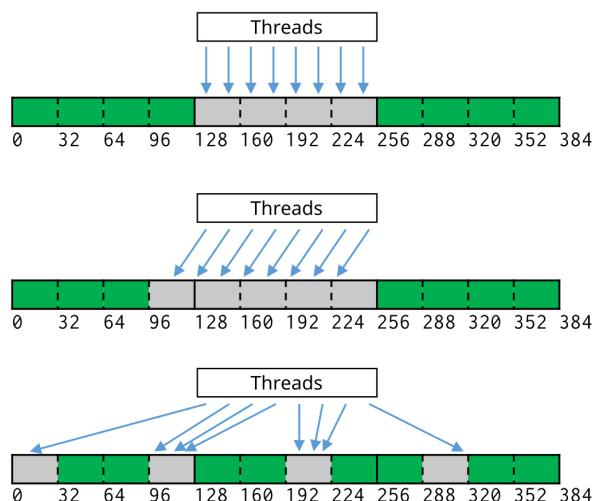
- Depends on the accelerator you are using
- You can try out different combinations, but deterministic optimizations require good knowledge on the accelerator hardware
  - In the case of NVIDIA GPUs you should start with the NVVP results and refer to CUDA documentation
  - One hard-coded value: for NVIDIA GPUs the vector length should always be 32, which is the (current) warp size

## Branches in device code

- 32 threads running the same instruction at the same time
- Avoid branches based on thread id unless evenly dividable by 32
  - If  $(i \% 2) \neq 0$
  - if  $(i \% 32) == 0$
- When unavoidable keep branches short

## Coalesced memory access

- Coalesced memory access
  - 32 threads accessing memory at the same time
  - 32 Byte access granularity
- Overly simplified
  - Some cases 128 bytes access granularity
  - 128 byte coalesced accesses can improve performance





## Summary

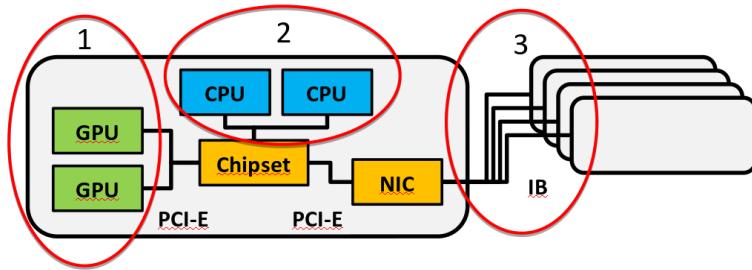
- Profiling is essential for optimization
  - NVPROF and NVVP for NVIDIA platform
- Loop optimizations
- Branches
- Memory access patterns

# OpenACC: multi-GPU programming



## Multi-GPU programming with OpenACC

- Three levels of hardware parallelism in a supercomputer:
  1. GPU - different levels of threads
  2. Node - multiple GPUs and CPUs
  3. System - multiple nodes connected with interconnect
- Three parallelization methods:
  1. OpenACC
  2. OpenMP or MPI
  3. MPI between nodes



## Multi-GPU communication cases

- Single node multi-GPU programming
  - All GPUs of a node are accessible from a single process and its OpenMP threads
  - Data copies either directly or through CPU memory
- Multi-node multi-GPU programming
  - Communication between nodes requires message passing (MPI)
- In this lecture we will discuss in detail only parallelization with MPI
  - It enables direct scalability from single to multi-node

## Multiple GPUs

- OpenACC permits using multiple GPUs within one node by using the `acc_get_num_devices` and `acc_set_device_num` functions
- Asynchronous OpenACC calls, OpenMP threads or MPI processes must be used in order to actually run kernels in parallel
- Issue when using MPI:
  - If a node has more than one GPU, all processes in the node can access all GPUs of the node
  - MPI processes do not have any a priori information about the other ranks in the same node
  - Which GPU the MPI process should select?

## Selecting a device with MPI

- Model is to use **one** MPI task per GPU
- Launching job
  - Launch your application so that there are as many MPI tasks per node as there are GPUs
  - Make sure the affinity is correct - processes equally split between the two sockets (that nodes typically have)
  - Read the user guide of the system for details how to do this!
- In the code a portable and robust solution is to use MPI3 shared memory communicators to split the GPUs between processes
- Note that you can also use OpenMP to utilize all cores in the node for computations on CPU side

## Selecting a device with MPI

```
MPI_Comm shared;
int local_rank, local_size, num_gpus;

MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, 0,
                     MPI_INFO_NULL, &shared);
MPI_Comm_size(shared, &local_size); // number of ranks in this node
MPI_Comm_rank(shared, &local_rank); // my local rank
num_gpus = acc_get_num_device(acc_device_nvidia); // num of gpus in node
if (num_gpus == local_size) {
    acc_set_device_num(local_rank);
} // otherwise error
```

## Data transfers

- Idea: use MPI to transfer data between GPUs, use OpenACC-kernels for computations
- Additional complexity: GPU memory is separate from CPU memory
- GPU-aware MPI-library
  - Can use the device pointer in MPI calls - no need for additional buffers
  - No need for extra buffers and device-host-device copies
  - If enabled on system, data will be transferred via transparent RDMA
- Without GPU-aware MPI-library
  - Data must be transferred from the device memory to the host memory and vice versa before performing MPI-calls

## Using device addresses with host\_data

- For accessing device addresses of data on the host OpenACC includes host\_data construct with the use\_device clause
- No additional data transfers needed between the host and the device, data automatically accessed from the device memory via **Remote Direct Memory Access**
- Requires *library* and *device* support to function!

## MPI communication with GPU-aware MPI

- MPI send
  - Send the data from the buffer on the **device** with MPI
- MPI receive
  - Receive the data to a buffer on the **device** with MPI
- No additional buffers or data transfers needed to perform communication

## MPI communication with GPU-aware MPI

```
/* MPI_Send with GPU-aware MPI */
#pragma acc host_data use_device(data)
{
    MPI_Send(data, N, MPI_DOUBLE, to, MPI_ANY_TAG, MPI_COMM_WORLD);
}

/* MPI_Recv with GPU-aware MPI */
#pragma acc host_data use_device(data)
{
    MPI_Recv(data, N, MPI_DOUBLE, from, MPI_ANY_TAG, MPI_COMM_WORLD,
              MPI_STATUS_IGNORE);
}
```

## Summary

- Typical HPC cluster node has several GPUs in each node
  - Selecting the GPUs with correct affinity
- Data transfers using MPI
  - GPU-aware MPI avoids extra memory copies